

# Markov Chain Monte Carlo

## How to Obtain Posterior

- Posterior distribution is proportional to a **product** of two distributions (prior and likelihood)
  - Conjugate – we know the class of the posterior distribution
  - Not conjugate – we don't know the distribution, simulation methods have to be adopted, Markov Chain Monte Carlo

## Markov Chain

- Markov Chain
  - Start with an initial point  $\theta_0$
  - The next point only relies on the value of the current point, using conditional probability  $p(\theta_{r+1}|\theta_r)$ , not the earlier stages
  - A nice property of Markov Chain is that, under certain conditions, the distribution of  $(\theta_r|\theta_0)$  will converge to a **stationary/invariant/equilibrium** distribution.
    - This distribution will have no influence by the value of  $\theta_0$ .

## Markov Chain Monte Carlo

- Markov Chain Monte Carlo
  - We want to establish the transition matrix  $p(\theta_{r+1}|\theta_r)$ , so that the simulated chain has a stationary distribution that is the posterior distribution we want.
  - The chain is **Ergodic** (遍历性), if it satisfies:

$$\lim_{R \rightarrow \infty} \frac{1}{R} \sum h(\theta_r) = E[h(\theta)]$$

- Law of large numbers
  - Standard version requires independence of the draws
  - As long as the dependence is not too strong, it still works

## Example of Markov Chain

Consider the Discrete case

$$S = \{q^1, q^2, \dots, q^d\}$$

d possible states –  
such as a  
discretization of  
parameter space

define a sequence of r.v.s given some initial starting point.

$$\Pr[\theta_{r+1} = \theta^j | \theta_r = \theta^i] = p_{i,j}$$

prob chain is in state j  
at “time” or iteration  
r+1 given in state i at  
iteration r

$$P = [p_{i,j}]; \sum_j p_{i,j} = 1$$

the rows of P specify the  
conditional distribution of  $\theta_{r+1}$   
given realized value of  $\theta_r$

5

## Stationary Distribution

Suppose we specify a distribution for  $\theta_0$ ,  $\pi_0$ . What is the distribution of  $\theta_1$ ?

$$\begin{aligned} \Pr[q_1 = q^j] &= \sum_i \Pr[q_0 = q^i] \Pr[q_1 = q^j | q_0 = q^i] \\ &= \sum_i p_{0,i} p_{i,j} \end{aligned}$$

or

$$p_1 = p_0 P, p_2 = p_1 P = p_0 P^2, \dots$$

row vectors

As r increases, the effects of starting  
distribution should “wear” off.

## Stationary Distribution

If we start in the stationary distribution, we should stay there which is why the stationary distribution is sometimes called the **invariant** distribution

$$\pi P = \pi \quad \text{eigenvector corresponding to eigenvalue of 1.}$$

$$\text{Let } P = \begin{bmatrix} 1/2 & 1/2 \\ 1/4 & 3/4 \end{bmatrix} \quad \pi = [1/3 \quad 2/3]$$

$$\pi P = [1/3 \quad 2/3] \begin{bmatrix} 1/2 & 1/2 \\ 1/4 & 3/4 \end{bmatrix} = [1/3 \quad 2/3]$$

## Stationary Distribution

If  $\pi_{i,j} > 0$ , then chain is called **irreducible** and there will be a stationary distribution

$$\lim_{r \rightarrow \infty} \pi_0 P^r = \pi$$

An irreducible chain is free to navigate. That is, we can get from anywhere to anywhere else (not necessarily in one step, though).

Ex:

$$P = \begin{bmatrix} .5 & .1 & .4 \\ 0 & .5 & .5 \\ 0 & .4 & .6 \end{bmatrix}$$

“trapped” in  
states 2 and 3

## Practical Considerations

- Effect of initial conditions
  - “burn-in” – run for B iterations, discard and use only the last R-B
- Non-iid Simulate – is this a problem?
  - No, Law of Large Numbers works for dependent sequences
  - Yes, simulation error larger than IID sequence
- How to construct the chain?

## First Method: Gibbs Sampler

- Posterior distribution for model parameters could be high-dimensional, which is hard to be simulated from
- Gibbs sampler provides a method to draw from the **conditionals** in order to achieve the joint.
- An example: to simulate the joint distribution of  $(\theta_1, \theta_2)$ 
  - Start with  $(\theta_1^0, \theta_2^0)$
  - Draw the next  $\theta_1^{r+1}$  from the conditional distribution  $\theta_1 | \theta_2^r$
  - Draw the next  $\theta_2^{r+1}$  from the conditional distribution  $\theta_2 | \theta_1^{r+1}$

## Why does it work?

- Suppose  $(\theta_1^r, \theta_2^r) \sim \pi(\cdot)$ , we need to prove that  $(\theta_1^{r+1}, \theta_2^{r+1}) \sim \pi(\cdot)$
- $(\theta_1^r, \theta_2^r) \sim \pi(\cdot)$ , therefore  $\theta_1^r \sim \pi_1(\cdot)$ , marginal distribution
- $\theta_2^{r+1} \sim \pi_{2|1}(\theta_2 | \theta_1^r)$  in Gibbs sampler
- Combine the last two steps, we get  $\theta_2^{r+1} \sim \int \pi_{2|1}(\theta_2 | \theta_1) \pi_1(\theta_1) d\theta_1 = \pi_2(\cdot)$
- Same to prove  $\theta_1^{r+1} \sim \pi_1(\cdot)$
- Therefore  $(\theta_1^{r+1}, \theta_2^{r+1}) \sim \pi(\cdot)$

## Gibbs Sampler

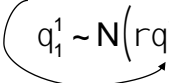
$$\theta \sim N\left(0, \begin{bmatrix} 1 & \rho \\ \rho & 1 \end{bmatrix}\right)$$

$$\theta_1 | (\theta_2 = a) \sim N\left(\mu_1 + \frac{\sigma_{12}}{\sigma_{22}}(a - \mu_2), \sigma_{11} - \frac{\sigma_{12}^2}{\sigma_{22}}\right)$$

$$q_2 | q_1 \sim N(rq_1, (1 - r^2)) \quad q_1 | q_2 \sim N(rq_2, (1 - r^2))$$

A simulator: Start at point  $\theta^0 = \begin{bmatrix} \theta_1^0 \\ \theta_2^0 \end{bmatrix}$

Draw  $\theta^1$  in two steps:

$$\begin{aligned} q_2^1 &\sim N(rq_1^0, 1 - r^2) \\ q_1^1 &\sim N(rq_2^1, 1 - r^2) \end{aligned}$$


Note: this is a Markov Chain. Current point entirely summarizes past.

# Gibbs Sampler

A simulator:


Start at point  $\theta^0 = \begin{bmatrix} \theta_1^0 \\ \theta_2^0 \end{bmatrix}$

Draw  $\theta^1$  in two steps:

$$\theta_2^1 \sim N(\rho\theta_1^0, 1-\rho^2)$$

$$\left( \theta_1^1 \sim N(\rho\theta_2^1, 1-\rho^2) \right)$$

repeat!



# General Gibbs Sampler

$$\theta = (\theta_1, \theta_2, \dots, \theta_p), \text{ "blocking"}$$

Sample from

$$\theta_{t+1,1} \sim f_1(\theta_1 | \theta_{t,2}, \theta_{t,3}, \dots, \theta_{t,p})$$

$$\theta_{t+1,2} \sim f_2(\theta_2 | \theta_{t+1,1}, \theta_{t,3}, \dots, \theta_{t,p})$$

$$\theta_{t+1,3} \sim f_2(\theta_3 | \theta_{t+1,1}, \theta_{t+1,2}, \dots, \theta_{t,p})$$

...

$$\theta_{t+1,p} \sim f_2(\theta_2 | \theta_{t+1,1}, \theta_{t+1,2}, \dots, \theta_{t+1,p-1})$$

# Applications of Gibbs Sampling

- Normal model
- Regression model

## Normal Model

- A set of random variable from normal distribution, knowing these variables  $x_1, x_2, \dots, x_N$ , use conjugate prior to calculate its mean  $\mu$ , and variance  $\sigma^2$

Likelihood function of the data, from the normal PDF

$$L(x|\mu, \sigma^2) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2} \frac{(x_i - \mu)^2}{\sigma^2}\right)$$

$$= \left(\frac{1}{\sqrt{2\pi\sigma^2}}\right)^N \exp\left(\sum_{i=1}^N -\frac{1}{2} \frac{(x_i - \mu)^2}{\sigma^2}\right)$$

Prior distribution for  $\mu \sim N(\mu_0, \sigma_0^2)$

Prior distribution for  $\sigma$

$$p(\sigma^2) \propto (\sigma^2)^{-\left(\frac{\nu_0}{2}+1\right)} \exp\left(-\frac{\nu_0 s_0^2}{2\sigma^2}\right) \sim \frac{\nu_0 s_0^2}{\chi_{\nu_0}^2}$$

Joint prior distribution  $p(\mu, \sigma^2) = p(\mu) \times p(\sigma^2)$



Posterior for  $\mu|\sigma^2 \sim N(\mu_1, \sigma_\mu^2)$

$$\mu_1 = \frac{\frac{\sum_1^N x_i}{\sigma^2} + \frac{\mu_0}{\sigma_0^2}}{\frac{N}{\sigma^2} + \frac{1}{\sigma_0^2}}, \quad \sigma_\mu^2 = \frac{1}{\frac{N}{\sigma^2} + \frac{1}{\sigma_0^2}}$$

Posterior for  $\sigma^2|\mu \sim \frac{\nu_1 s_1^2}{\chi^2(\nu_1, s_1^2)}$ , also called *scaled - inv -  $\chi^2$* ( $\nu_1, s_1^2$ ),  
or *Inv - Gamma*  $\left(\frac{\nu_1}{2}, \frac{\nu_1 s_1^2}{2}\right)$

$$\nu_1 = \nu_0 + N, \quad s_1^2 = \frac{\nu_0 s_0^2 + N s^2}{\nu_0 + N}$$

## Regression Model

- Regression model

$$y = X\beta + \epsilon, \epsilon \sim N(0, \sigma^2)$$

- Likelihood function

$$L(y, X|\beta, \sigma^2) = l(\beta, \sigma^2) \\ \propto \frac{1}{2\pi\sigma^N} \exp\left(-\frac{1}{2} \sum_i \frac{(y_i - X_i\beta)^2}{\sigma^2}\right)$$

- Prior distribution

$$p(\beta, \sigma^2) = p(\beta)p(\sigma^2)$$

## Bayes Regression

- Prior probabilities

$$p(\beta) \propto \exp\left[\frac{-1}{2}(\beta - \bar{\beta})' A(\beta - \bar{\beta})\right]$$

$$p(\sigma^2) \propto (\sigma^2)^{-\left(\frac{\nu_0}{2} + 1\right)} \exp\left[\frac{-\nu_0 S_0^2}{2\sigma^2}\right]$$

19

- Posterior of all the model parameters  $\beta, \sigma^2$

$$\begin{aligned} & l(\beta, \sigma^2) \times \text{prior}(\beta, \sigma^2) \\ &= \frac{1}{2\pi\sigma} \exp\left(-\frac{1}{2\sigma^2} \sum_i (y_i - X_i\beta)^2\right) \\ &\times \exp\left(-\frac{1}{2}(\beta - \bar{\beta})' A(\beta - \bar{\beta})\right) \\ &\times (\sigma^2)^{-\left(\frac{\nu_0}{2} + 1\right)} \exp\left(\frac{-\nu_0 S_0^2}{2\sigma^2}\right) \end{aligned}$$

## Gibbs Sampler

Scheme:  $[y|X, \beta, \sigma^2] [\beta] [\sigma^2]$

- 1) Draw  $[\beta | y, X, \sigma^2]$
- 2) Draw  $[\sigma^2 | y, X, \beta]$  (conditional on  $\beta$ !)
- 3) Repeat

21

## Gibbs Sampler

- Full Conditional distribution of  $\beta$ , and  $\sigma^2$

$$[b | y, X, s^2] = N(\tilde{b}, (s^{-2}X'X + A)^{-1})$$

with  $\tilde{b} = (s^{-2}X'X + A)^{-1}(s^{-2}X'X\hat{b} + A\bar{b})$   
 $\hat{b} = (X'X)^{-1}X'y$

$$[s^2 | y, X, b] = \frac{n_1 s_1^2}{C_{n_1}} \text{ with } n_1 = n_0 + n$$

$$s_1^2 = \frac{n_0 s_0^2 + (y - X\bar{b})'(y - X\bar{b})}{n_0 + n}$$

Depends on  $\beta$

22

## Summary

- MCMC basic concepts
- First MCMC: Gibbs sampler
  - Gibbs sampler and Normal model
  - Gibbs sampler for Normal regression
- Second MCMC: Metropolis-Hastings
  - Next class